

Research Statement

Hanno Foerster

My research studies how labor market outcomes emerge from the joint decision-making of family members, with a particular focus on frictions that shape family labor supply, migration, and household formation. The growing availability of large-scale administrative data that link family members provides new empirical opportunities to study decisions made within families. I combine empirical analysis with structural modeling to uncover the mechanisms underlying these decisions.

The main goal of my research agenda is to advance our understanding of how families interact with labor markets and to generate policy-relevant insights for labor market and family policy. Examples of topics I study include the economic effects of divorce law changes, the interaction between income inequality and household formation, the political economy of labor laws, and the role of remote work in shaping household labor market outcomes.

In addition to my primary line of research on families and labor markets, my work comprises a secondary strand that examines how labor market policies affect job search behavior and matching outcomes. I summarize both strands of my research agenda below.¹

Family Decisions and Economic Policy

Economic policies shape family formation, marital stability, and labor supply within households. Understanding how policies, including tax and transfer systems, divorce law, and labor regulation, influence these family decisions is important for evaluating their broader economic consequences. For example, divorce is a common cause of poverty (especially for women), marital sorting shapes the distribution of income and wealth across households, and gender differences in labor supply contribute to persistent gender income inequality.

One important policy margin is divorce law, which commonly mandates financial transfers between ex-spouses, providing insurance against the economic consequences of divorce. In my paper [Untying the Knot: How Child Support and Alimony Affect Couples' Dynamic Decisions and Welfare](#) ([1], solo authored), I study how child support and alimony policies can balance competing policy objectives. In particular, I ask how such policies can balance the trade-off between providing insurance to the lower earner and enabling efficient specialization within couples, while limiting distortions to post-divorce labor supply incentives.

Prior work has generally studied child support and alimony through the lens of static models. I analyze this policy trade-off in a dynamic framework that accounts for married couples' risk of future divorce and the resulting weaker incentives for household specialization. The framework allows for self-insurance through the accumulation of assets and human capital, and captures dynamic adjustments in labor supply, savings, and consumption around divorce.

I apply the model to Denmark, where child support and alimony payments are governed by well-defined institutional rules that are relatively strictly enforced. I draw on administrative records linking marital histories, labor market careers, tax and birth records, and the exact amounts of child support and alimony paid and received after divorce. The model is estimated by matching a broad set of empirical patterns, including household specialization during marriage and adjustments around divorce.

I find that within-couple gender inequality in consumption increases sharply following divorce and remains persistently elevated. This pattern arises despite self-insurance through savings and human capital accumulation and despite existing child support and alimony policies. The result is tightly grounded in the data: consumption can be backed out from observed income and assets under a small set of structural assumptions, and is robust to accounting for repartnering and partners' income and assets.

¹Up-to-date versions of my CV, papers, and their statuses, are available on my website: www.hannofoerster.com.

Model simulations show that alimony and child support can mitigate this increase in consumption inequality following divorce, although their effects are dampened by distortions to labor supply. The model highlights two key frictions that create scope for policy: limited commitment within marriage (partners cannot perfectly commit to future allocations) and non-cooperation after divorce. These frictions amplify the decline in the lower-earner's consumption following divorce and reduce post-divorce labor supply incentives. They thereby generate welfare losses that are partially offset by child support and alimony policies. Among the counterfactual policies considered, those that come closest to the efficient frontier are backward-looking, relying only on economic outcomes realized before divorce and thereby eliminating strategic responses in ex-spouses' labor supply decisions. This paper was published in the *Review of Economic Studies* in 2025.

Marital sorting is a key channel through which household formation affects aggregate economic outcomes. In the paper [Job Displacement, Remarriage, and Marital Sorting](#) ([2], joint with Tim Obermeier and Bastian Schulz), we study how adverse labor market shocks affect household formation, dissolution, and marital sorting.

Previous work has studied marital sorting by examining cross-sectional correlations in spouses' characteristics. We instead compare individuals who experience job displacement with a non-displaced control group and follow their transitions into and out of marriages over time. This design yields novel within-person evidence on how marital sorting responds to quasi-exogenous labor market shocks, offering new insights into the mechanisms shaping marital sorting.

We study these dynamics in Denmark using administrative records that track individuals' labor market histories and their transitions into and out of marriages and cohabiting relationships. We find that job displacement persistently reduces labor income, lowers marital stability, and increases the probability of remaining single. Displaced individuals, on average, transition away from lower-earning partners and toward higher-earning partners relative to the control group. This pattern of upward rematching following job displacement is difficult to reconcile with the positive association in spouses' earnings in cross-sectional data.

To interpret these patterns, we use a marriage market search and matching framework that builds on the Shimer–Smith model. We show formally that the standard interpretation of marital sorting as reflecting positive assortative matching cannot jointly explain the upward rematching observed following job displacement and the positive cross-sectional association in spouses' earnings.

We propose multidimensional matching as an explanation that reconciles both patterns. If partners' incomes enter marital surplus as substitutes, as in Becker's classic theory of marriage, sorting on income becomes negative, while sorting on other characteristics may remain positive. When these characteristics are positively correlated with income, the model can generate both the upward rematching following job displacement and the positive cross-sectional association in spouses' earnings.

The multidimensional framework also allows us to quantify the sources of the income correlation between spouses. Our estimates show that sorting on age and education accounts for at least 56% of this correlation, with the remainder explained by other factors, including unobserved characteristics.

To illustrate the implications for counterfactual analysis, we simulate an increase in individual income inequality and examine the implications for household-level inequality. Under one-dimensional positive assortative matching, marital sorting amplifies the increase in household-level inequality, as high-income individuals are more likely to match with other high-income individuals. In contrast, in the multidimensional model calibrated to match both our job-displacement evidence and the cross-sectional income correlation, it dampens this increase, as sorting on income conditional on other characteristics is negative. This paper is currently under revision for resubmission to the *American Economic Journal: Macroeconomics*.

Historically, advances in women's economic rights have been central to the evolution of women's role in the labor market. In the paper [Protection for Whom? The Political Economy of Protective Labor Laws for Women](#)

([3], joint with Matthias Doepke, Anne Hannusch, and Michèle Tertilt), we study the political economy forces that gave rise to gender-specific labor laws in the late nineteenth and early twentieth centuries. These laws were framed as protections for women but restricted their employment in several ways, including limits on working hours, bans on night work, and workplace requirements such as mandatory seating provisions and limits on heavy lifting.

Previous economic research has studied the consequences of protective legislation, while historians have documented the stated motivations of proponents and opponents of gender-specific labor laws. We instead analyze economic mechanisms that contributed to the broad political support these laws gained. This support drove their widespread adoption across U.S. states in the early twentieth century and, decades later, their repeal.

Specifically, we examine whether changes in labor market competition between men and women and in the composition of voters affected by it can explain the rise and decline of protective labor laws. To formalize this mechanism, we develop a political economy model with an agricultural and a modern sector. Households choose labor supply and sector of employment and vote on the introduction or repeal of gender-specific restrictions. If enacted, these restrictions lower women's productivity in the modern sector, which in equilibrium reduces their wages while increasing the wages of low-skilled men who compete with them. The resulting changes in household income, together with adjustments in goods prices, shape the voting incentives of different demographic groups. The model predicts that political coalitions are not simply divided along gender lines, but instead depend on marital status, sector of employment, and men's skill level.

To test the model's predictions, we assemble a novel dataset from historical sources covering the enactment of five types of protective labor laws across U.S. states. Our results show that, at the aggregate level, the relative size of groups predicted to support protective laws closely tracks both the rise of protective legislation between 1910 and 1920 and its decline in the 1960s. Exploiting variation across states and over time, we further find that changes in the prevalence of these groups systematically coincide with the introduction and repeal of protective legislation at the state level (controlling for time and state fixed effects).

To identify the key drivers of shifts in support for protective legislation, we calibrate the model and conduct counterfactual simulations. We find that the emergence of the modern sector, where male and female workers compete more strongly than in agriculture, is the main driver of the rise of protective labor laws. The rise in married women's labor force participation after World War II, which increasingly aligns the economic interests of women and men, is the key driver behind the repeal of protective labor laws. This paper is a reject and resubmit at the *American Economic Review*.

Dual-Earner Households and Spatial Frictions

Given the high and rising share of dual-earner households, understanding how these households navigate labor markets is increasingly important. Labor market decisions within dual-earner households are subject to unique constraints and frictions, including the spatial constraint that they must secure two jobs within the same commuting area. This constraint shapes their job search and labor market outcomes and introduces additional frictions in the allocation of workers to jobs.

In [The Colocation Friction: Dual-Earner Job Search, Migration, and Labor Market Outcomes](#) ([4], joint with Robert Ulbricht), we develop a spatial search framework to formalize, characterize, and quantify how search frictions shape the job search, migration, and labor market outcomes of dual-earner households.

Prior work has shown that dual-earner households are less likely to migrate for work, that migration often leaves one spouse temporarily unemployed, and that women are more likely to take on this trailing-spouse role. We formalize and study a key friction underlying these patterns: dual-earner households require a double coincidence of job offers to relocate without one spouse becoming unemployed, yet employers do not coordinate to generate such coincidences. A key feature of our framework, relative to the previous literature, is

that search is directed, allowing spouses to optimally coordinate their job search across locations in response to this friction.

We characterize the friction by showing formally when it leads to a strict welfare loss relative to a benchmark economy that captures employer coordination. In the benchmark economy, the feasibility frontier for total matches is held fixed, but matches are redistributed across households to generate double coincidences of job offers where feasible and welfare-improving. We derive a general theorem characterizing when the benchmark strictly dominates the baseline economy and identify key forces that amplify the welfare loss from the friction, including (i) high migration costs, (ii) high returns to work experience that increase the cost of migrating with a trailing spouse, and (iii) search efficiency losses from narrowing job search to a single location.

To study the consequences of the friction, we calibrate the model to the U.S. at the commuting-zone level, drawing on a variety of micro datasets to inform different model components. The calibrated model captures distances between commuting zones, distance-dependent migration rates, and geographic variation in wages and amenities. Amenities are informed by data including local school quality, crime rates, public expenditures, and weather conditions.

The estimated model replicates the empirical pattern that, following migration, within-couple earnings and wage gaps widen as women are more likely than men to experience unemployment. The model predicts that women assume the trailing-spouse role in roughly 80% of migrating households. A decomposition reveals that this asymmetry is driven primarily by gender differences in average productivity (capturing a range of factors underlying the gender wage gap).

Model simulations show that the friction substantially reduces women's long-term earnings gains from migration by 19% and accounts for 18% of the gender gap in migration gains. It lowers migration rates by 13%, with effects concentrated among high-income dual-earner households ("power couples").

We simulate the rise of remote work by allowing jobs to convert to remote positions at rates calibrated to occupation-level shares of remote-workable jobs. Migration initially declines as households delay relocation due to the option value of waiting for remote conversion, before recovering gradually as more jobs convert. In the long run, remote work halves the earnings losses caused by the friction. This paper has been submitted.

I am also working on a project with Felix Koenig and Jennifer Lee Bernard that studies how families navigate colocation constraints. Using U.S. employer-employee linked data matched to tax records, we exploit establishment relocations that quasi-exogenously shift workers' jobs across commuting zones. This setting allows us to examine how changes in job location affect migration behavior, employment, and earnings for both workers and their spouses. The project aims to provide new evidence on how colocation constraints shape gender differences in labor market outcomes.

Labor Market Policies and Job Search Behavior

In addition to studying family decision-making, my research examines how active labor market policies affect job search behavior and equilibrium labor market outcomes. This work combines the analysis of rich administrative and survey data with structural modeling to quantify the responses of workers and firms to both existing and counterfactual policies.

A common feature of active labor market policies is that benefit recipients are required to apply to jobs referred by caseworkers and may face sanctions if they refuse these vacancy referrals. However, monitoring of compliance is imperfect, and workers may exploit institutional features of the system to avoid their obligations. In [Structural Empirical Analysis of Vacancy Referrals with Imperfect Monitoring and the Strategic Use of Sickness Absence](#) ([5], joint with Gerard van den Berg and Arne Uhlenborff), we analyze one channel of avoidance that has received little attention in the previous literature: strategic sick reporting. In the institutional setting we study, unemployed workers who report sick are temporarily exempt from search

requirements and can thereby avoid sanctions. We examine how strategic sick reporting affects job search behavior and the effectiveness of labor market policy.

We develop a structural job search model that incorporates the key features of the German monitoring system and estimate it using rich German administrative data. The data link labor market biographies from social security records with detailed caseworker-recorded information on vacancy referrals, sanctions, and sick reporting. In the model, workers optimally decide whether to report sick after receiving a vacancy referral. Their strategy is characterized by three reservation wages: one for regular (non-referral) offers, one that triggers strategic sick reporting following a referral, and one for rejecting a referral with the risk of a sanction. We show that all model parameters are identified given our data. In the structural estimation, we allow for rich observed and unobserved heterogeneity across demographic groups.

Our estimates show that about 9% of sick reports among unemployed workers are induced by vacancy referrals. Counterfactual policy simulations show that stronger sanction enforcement and higher vacancy referral rates both have the potential to substantially reduce mean unemployment duration. The two policies complement each other, with stronger effects when implemented jointly. Strategic sickness reporting reduces the effectiveness of either policy by about 7% and that of joint implementation by about 3%. This paper has been accepted for publication in the *Journal of Labor Economics*.

One key motivation for using vacancy referrals in active labor market policy is to bring together unemployed workers and firms that otherwise would not have matched. Existing empirical evidence shows that vacancy referrals significantly increase the job-finding probabilities of recipients. In [The Equilibrium Effects of Vacancy Referrals](#) ([6], joint with Gerard van den Berg), we examine to what extent these positive individual-level effects are attenuated by equilibrium effects. To address this question, we develop a search-and-matching framework that incorporates three channels through which equilibrium effects may arise: (i) congestion in firms' applicant pools, whereby referred workers may crowd out other applicants, (ii) disincentive effects on workers' search effort (both from receiving referrals and from increased congestion), and (iii) endogenous changes in firms' vacancy posting. In the model, unemployed workers choose how many applications to submit and may receive vacancy referrals. Matching follows an urn-ball technology. Firms post vacancies and wages, and select one worker from the resulting applicant pool.

We estimate the model using worker and firm survey data from Germany. The worker data provide detailed information on unemployed individuals' job search behavior, including the number of applications submitted, the number of vacancy referrals received, and post-unemployment outcomes. The firm data contain vacancy-level information on applicant pools and hiring outcomes, including the number of applicants and the number who arrive via vacancy referrals.

Counterfactual policy simulations show that equilibrium effects substantially reduce the effectiveness of vacancy referrals. A large-scale reform that increases the referral rate by about 60% would raise the worker match rate by 0.9 percentage points. Absent equilibrium effects, the effect would be about 30% larger. We further decompose these effects and show that equilibrium forces reduce worker match rates by similar magnitudes for newly treated individuals, individuals already treated prior to the reform, and untreated individuals. We are currently finalizing this paper for submission.

References

1. Foerster, H. (2025), Untying the Knot: How Child Support and Alimony Affect Couples' Dynamic Decisions and Welfare, *The Review of Economic Studies* Volume 92, Issue 5, October 2025, Pages 3029–3066.
2. Foerster, H., Obermeier, T., and Schulz, B. (2025), Job Displacement, Remarriage, and Marital Sorting, *American Economic Journal: Macroeconomics*, revise and resubmit.

3. Doepke, M., Foerster, H., Hannusch, A., and Tertilt, M. (2025), Protection for Whom? The Political Economy of Protective Labor Laws for Women, *American Economic Review*, reject and resubmit.
4. Foerster, H., and Ulbricht, R. (2025), The Colocation Friction: Dual-Earner Job Search, Migration, and Labor Market Outcomes, *Boston College Working Paper No. 1103*.
5. van den Berg, G., Foerster, H., and Uhlendorff, A. (2025), Structural Empirical Analysis of Vacancy Referrals with Imperfect Monitoring and the Strategic Use of Sickness Absence, *Journal of Labor Economics*, forthcoming.
6. van den Berg, G., and Foerster, H. (2025), The Equilibrium Effects of Vacancy Referrals, *Working Paper*.